

Statistics with R - Exercise 1

Philipp Satlawski - h0640348

06/11/2020

This document contains the answered questions of exercise 1 for the course “Statistics with R”.

Task 1 - Sequences

Creation and comparison of vectors.

1. Create a vector x that contains the sequence of even numbers from 0 to 40 ($x \in \mathbb{R}$)

```
(x <- seq(from = 0, to = 40, by = 2))
```

```
## [1] 0 2 4 6 8 10 12 14 16 18 20 22 24 26 28 30 32 34 36 38 40
```

2. Create a vector y , which contains the elements of vector x but in *random* order ($y \in x$)

```
# set seed to produce reproducible results  
set.seed(10)  
# take random samples from the elements of vector x  
(y <- sample(x))
```

```
## [1] 20 16 18 30 22 12 14 10 34 4 28 36 24 2 40 8 32 38 6 0 26
```

3. The values of x and y agree on the following positions:

```
# number of positions the two vectors agree on  
length(which(x == y))
```

```
## [1] 2
```

```
# indices (positions) on which the values of the two vectors agree on  
which(x == y)
```

```
## [1] 13 17
```

Task 2 - Sequences

Verification of three approximation formulas for π . Since the calculation of π by the following formulas includes infinite sums or products, we use a large $n = 10000$.

Creating vector `i` as a sequence from 1 to 100000 ($x \in \mathbb{R}$).

```
# create vector i  
i <- 1:10000
```

1. John Wallis (1616-1703):

```
# calculation of pi with the formula of John Wallis  
prod((2*i/(2*i-1)) * (2*i/(2*i+1)))*2
```

```
## [1] 3.141514
```

2. Gottfried Leibnitz (1646-1716):

```
# calculation of pi with the formula of Gottfried Leibnitz  
sum((( -1)**(i+1))/(2*i-1))*4
```

```
## [1] 3.141493
```

3. Leonhard Euler (1707-1783):

```
# calculation of pi with the formula of Leonhard Euler  
sqrt(sum(1/i**2)*6)
```

```
## [1] 3.141497
```

To compute the smallest relative deviation from π given the vector `i`, we need π 's current best approximation. According to [ikipedia](#) (2020) π 's current best approximation is 3.1415926535897932384626433.

```
# store pi's current best approximation for comparison  
pi <- 3.1415926535897932384626433
```

```
# calculate the difference between John Wallis (1616-1703) and current pi  
(diff_JW <- abs(prod((2*i/(2*i-1)) * (2*i/(2*i+1)))*2 - pi))
```

```
## [1] 7.853491e-05
```

```
# calculate the difference between Gottfried Leibnitz (1646-1716) and current pi  
(diff_GL <- abs(sum((( -1)**(i+1))/(2*i-1))*4 - pi))
```

```
## [1] 1e-04
```

```
# calculate the difference between Leonhard Euler (1707-1783) and current pi
(diff_LE <- abs(sqrt(sum(1/i**2)*6) - pi))
```

```
## [1] 9.548964e-05
```

```
# get the absolute difference of the formula with the smallest absolute difference
min(diff_JW, diff_GL, diff_LE)
```

```
## [1] 7.853491e-05
```

Hence, the approximation formula for π with the smallest relative deviation for $n = 10000$ is John Wallis (1616-1703) formula.

Wikipedia (2020), https://en.wikipedia.org/wiki/Approximations_of_%CF%80

Task 3 - Vectors

Creation and manipulation of vectors.

1. Create vector x containing the sequence 1 to 100 ($x \in \mathbb{R}$) and vector 'y' containing a sample of size $n = 70$ of the sequence 1 to 150 ($y \in \mathbb{R}$).

```
# create vector x containing the sequence 1 to 100 (natural numbers)
(x <- 1:100)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18
## [19] 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36
## [37] 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54
## [55] 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
## [73] 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90
## [91] 91 92 93 94 95 96 97 98 99 100
```

```
# set seed to produce reproducible results
set.seed(10)
# create vector y containing a sample of size n = 70 of the sequence 1 to 150 (natural numbers)
(y <- sample(1:150, 70, replace = TRUE))
```

```
## [1] 137 74 112 72 88 15 143 74 24 13 95 136 110 7 86 82 29 29 121
## [20] 92 50 109 101 122 33 135 68 93 114 88 51 32 11 79 92 91 42 78
## [39] 13 105 144 117 26 89 48 15 110 24 61 132 14 35 10 74 58 144 15
## [58] 31 138 101 101 109 39 118 89 18 131 42 138 79
```

2. Determine the amount of elements that are contained in x but not in y .

```
# elements of x not contained in y
(setdiff(x,y))
```

```
## [1] 1 2 3 4 5 6 8 9 12 16 17 19 20 21 22 23 25 27 28
## [20] 30 34 36 37 38 40 41 43 44 45 46 47 49 52 53 54 55 56 57
## [39] 59 60 62 63 64 65 66 67 69 70 71 73 75 76 77 80 81 83 84
## [58] 85 87 90 94 96 97 98 99 100
```

```
# number of elements of x not contained in y
length(setdiff(x,y))
```

```
## [1] 66
```

3. Check for duplicate elements in `y` and depending if there are duplicate elements or not create a different `z`.

```
# check if there are duplicate elements in y
if(length(y[duplicated(y)]) > 0) {
  # create a new vector z containing the duplicate elements of y
  z <- y[duplicated(y)]
} else {
  # create a new vector z as a copy of y
  z <- y
}
```

4. Determine the number of elements of `z` that are multiples of 3.

```
# calculate the number of elements of z that are multiples of 3
length(z[z %% 3 == 0])
```

```
## [1] 6
```

5. Revert the `y` without using the function `rev()`.

```
# revert vector y
y[length(y):1]
```

```
## [1] 79 138 42 131 18 89 118 39 109 101 101 138 31 15 144 58 74 10 35
## [20] 14 132 61 24 110 15 48 89 26 117 144 105 13 78 42 91 92 79 11
## [39] 32 51 88 114 93 68 135 33 122 101 109 50 92 121 29 29 82 86 7
## [58] 110 136 95 13 24 74 143 15 88 72 112 74 137
```

Task 4 - Point Estimation

Assuming a normally distributed population, we create random sample and estimate μ and σ^2 for this sample.

1. Draw a reproducible sample of size $n = 30$ from a normal distribution with $\mu = 5$ and $\sigma^2 = 4$.

```
# set seed to produce reproducible results
set.seed(10)
# draw random sample with n = 30, mu = 5, sigma^2 = 4
(x <- rnorm(30, mean = 5, sd = sqrt(4)))
```

```
## [1] 5.0374923 4.6314949 2.2573389 3.8016646 5.5890903 5.7795886 2.5838476
## [8] 4.2726480 1.7466546 4.4870432 7.2035590 6.5115630 4.5235329 6.9748894
## [15] 6.4827803 5.1786945 3.0901123 4.6096992 6.8510425 5.9659570 3.8073787
## [22] 0.6294263 3.6502681 0.7618776 2.4696040 4.2526769 3.6248891 3.2556823
## [29] 4.7964780 4.4924389
```

2. Estimate μ and σ^2 on the basis of your sample using the formulas to estimate the population mean μ with the sample mean \bar{x} and the population variance σ^2 with the empirical variance s^2 , without using the functions `mean()`, `var()` and `sd()`.

```
# calculation of mean
(sum(x)/length(x))
```

```
## [1] 4.310647
```

```
# calculation of variance
(1/(length(x)-1) * sum((x-mean(x))**2))
```

```
## [1] 3.00578
```

```
# calculation of standard deviation
(sqrt(1/(length(x)-1) * sum((x-mean(x))**2)))
```

```
## [1] 1.733718
```

3. Compare your results with the output of the functions `mean()` and `var()`.

```
# compute the mean with inbuilt function mean()
(mean(x))
```

```
## [1] 4.310647
```

```
# check if the results of the inbuilt function mean() and the formula for calculating the mean are the same
if(round(sum(x)/length(x),5) == round(mean(x),5)) {
  ("same")
} else {
  ("different")
}
```

```
## [1] "same"
```

```
# compute the variance with inbuilt function var()
(var(x))
```

```
## [1] 3.00578
```

```

# check if the results of the inbuilt function var() and the formula for calculating the variance are the same
if(round((1/(length(x)-1) * sum((x-mean(x))**2)),5) == round(var(x),5)) {
  ("same")
} else {
  ("different")
}

```

```
## [1] "same"
```

```

# compute the standard variation with inbuilt function sd()
(sd(x))

```

```
## [1] 1.733718
```

```

# check if the results of the inbuilt function sd() and the formula for calculating the variance are the same
if(round(sqrt(1/(length(x)-1) * sum((x-mean(x))**2)),5) == round(sd(x),5)) {
  ("same")
} else {
  ("different")
}

```

```
## [1] "same"
```

4. Are your estimates close to the population values? Repeat the steps 1 and 3 from above with a sample of size $n = 3000$. What do we learn?

```

# set seed to produce reproducible results
set.seed(10)
# draw random sample with n = 3000, mu = 5, sigma^2 = 4
x <- rnorm(3000, mean = 5, sd = sqrt(4))

# calculation of mean
(sum(x)/length(x))

```

```
## [1] 5.002634
```

```

# calculation of variance
(1/(length(x)-1) * sum((x-mean(x))**2))

```

```
## [1] 4.143197
```

```

# calculation of standard deviation
(sqrt(1/(length(x)-1) * sum((x-mean(x))**2)))

```

```
## [1] 2.035484
```

```

# compute the mean with inbuilt function mean()
(mean(x))

```

```
## [1] 5.002634
```

```
# check if the results of the inbuilt function mean() and the formula for calculating the mean are the same
if(round(sum(x)/length(x),5) == round(mean(x),5)) {
  ("same")
} else {
  ("different")
}
```

```
## [1] "same"
```

```
# compute the variance with inbuilt function var()
(var(x))
```

```
## [1] 4.143197
```

```
# check if the results of the inbuilt function var() and the formula for calculating the variance are the same
if(round((1/(length(x)-1) * sum((x-mean(x))**2)),5) == round(var(x),5)) {
  ("same")
} else {
  ("different")
}
```

```
## [1] "same"
```

```
# compute the standard variation with inbuilt function sd()
(sd(x))
```

```
## [1] 2.035484
```

```
# check if the results of the inbuilt function sd() and the formula for calculating the variance are the same
if(round(sqrt(1/(length(x)-1) * sum((x-mean(x))**2)),5) == round(sd(x),5)) {
  ("same")
} else {
  ("different")
}
```

```
## [1] "same"
```

The conclusion of the above calculations shows, the higher the sample size n the smaller is the deviation between the sample mean \bar{x} and the population mean μ . The same applies for the variance and standard deviation.

Task 5 - Interval Estimation

Calculate the confidence intervals for the mean and the variance.

1. Draw a reproducible sample of size $n = 30$ from a normal distribution with $\mu = 5$ and $\sigma^2 = 4$.

```
# set seed to produce reproducible results
set.seed(10)
# draw random sample
(x <- rnorm(30, mean = 5, sd = sqrt(4)))
```

```
## [1] 5.0374923 4.6314949 2.2573389 3.8016646 5.5890903 5.7795886 2.5838476
## [8] 4.2726480 1.7466546 4.4870432 7.2035590 6.5115630 4.5235329 6.9748894
## [15] 6.4827803 5.1786945 3.0901123 4.6096992 6.8510425 5.9659570 3.8073787
## [22] 0.6294263 3.6502681 0.7618776 2.4696040 4.2526769 3.6248891 3.2556823
## [29] 4.7964780 4.4924389
```

2. Calculate a confidence interval for μ and σ^2 for $\alpha = 0.05$ (hence the confidence level is $1 - \alpha = 0.95$). Inbuilt functions such as `mean()`, `sd()` and `var()` are allowed. looked up in skriptum $t(29;0.975) = 2.045$

```
# looked up the t-value in the skript t(29;0.975) = 2.045
# get t-value for n-1 = 29; 1-alpha = 0.975 from an inbuilt R function
(t <- qt(0.975,df=29))
```

```
## [1] 2.04523
```

```
# lower bound mean
(low_bound_mean <- mean(x) - t * sd(x)/sqrt(length(x)))
```

```
## [1] 3.663266
```

```
# upper bound mean
(up_bound_mean <- mean(x) + t * sd(x)/sqrt(length(x)))
```

```
## [1] 4.958028
```

```
# get chi-value for n-1 = 29; 1-alpha = 0.975
(qchisq(0.975, df=29))
```

```
## [1] 45.72229
```

```
# lower bound standard dev
(low_bound_var <- ((length(x)-1) * var(x)) / qchisq(0.975, df=29))
```

```
## [1] 1.906458
```

```
# upper bound standard dev
(up_bound_var <- ((length(x)-1) * var(x)) / qchisq(0.025, df=29))
```

```
## [1] 5.431995
```

3. Determine if true parameters lie in the confidence interval.


```

# check if the mean lies within the confidence interval of the mean
if((mean(x) > low_bound_mean) & (mean(x) < up_bound_mean)){
  ("the mean lies within the confidence interval")
} else{
  ("the mean lies outside the confidence interval")
}

```

```
## [1] "the mean lies within the confidence interval"
```

```

# check if the variance lies within the confidence interval of the variance
if((var(x) > low_bound_var) & (var(x) < up_bound_var)){
  ("the variation lies within the confidence interval")
} else{
  ("the variation lies outside the confidence interval")
}

```

```
## [1] "the variation lies within the confidence interval"
```

Yes, in our case the true parameters lie within our confidence intervals. This is true for both the mean and the variance. The mean and the variance are in within the confidence intervals with a 5% probability of error.